

Artificial intelligence in data science

Long Short-Term Memory Networks

Janos Török

Department of Theoretical Physics


October 20, 2022

Visualization of different layers


- ▶ Algorithm:
- ▶ Use the optimization algorithm of tensorflow and maximize the outcome of one class
- ▶ We end up with images which will be 100% in one category
- ▶ How they look like?

Visualization of different layers

conv1_1: a few of the 64 filters




conv2_1: a few of the 128 filters



Visualization of different layers

conv4_1: a few of the 512 filters



conv5_1: a few of the 512 filters




Visualization of different layers



Hartebeest



Measuring Cup



Ant



Starfish



Anemone Fish



Banana




Parachute




Screw


Sea snake




Rocking Chair



Rugby ball




Lake side



Cheat neural networks

- ▶ If neural network is exactly known
- ▶ Images can be made which are categorized falsely
- ▶ You can hide your activity




Temporal data

- ▶ Most of the data is sequential, can be ordered
- ▶ Very often time orders the data
- ▶ Prediction is very important
- ▶ For this we need history



Source: Akshay Sood

Feedforward neural network



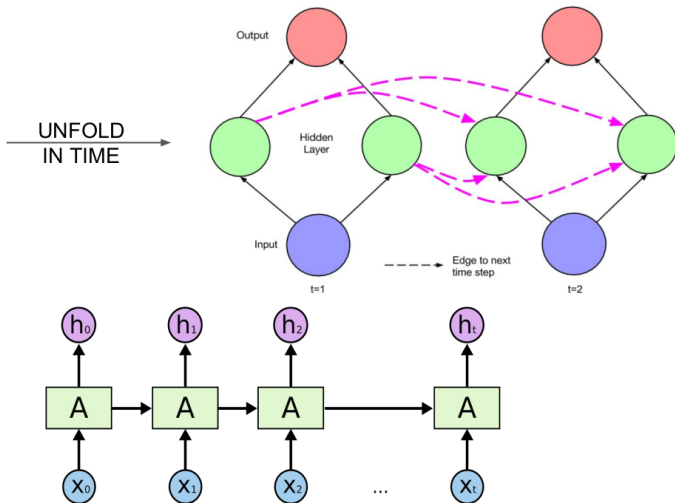
Recurrent Neural Networks (RNN)

- Output depends on previous state and current output




Recurrent Neural Networks (RNN)

- ▶ Output depends on previous state and current output
- ▶ Feedback loops




Training RNNs

- ▶ Backpropagation through time
- ▶ Regular (feedforward) backprop applied to RNN unfolded in time




Training RNNs

- ▶ Backpropagation through time
- ▶ Regular (feedforward) backprop applied to RNN unfolded in time
- ▶ Problem: can't capture long-term dependencies due to vanishing/exploding gradients during backpropagation



Training RNNs

- Problem: can't capture long-term dependencies due to vanishing/exploding gradients during backpropagation




$$h^{(t)} = \sigma(w_c \cdot c^{(t)})$$

$$c^{(t)} = \sigma(w_r \cdot c^{(t-1)} + w_x \cdot x^{(t)})$$

Long Short-Term Memory networks (LSTM)


- ▶ A type of RNN architecture that addresses the vanishing/exploding gradient problem and allows learning of long-term dependencies
- ▶ Recently risen to prominence with state-of-the-art performance in speech recognition, language modeling, translation, image captioning




LSTM Memory Cell

Long Short-Term Memory networks (LSTM)

- ▶ Memory cell (block): maintains its state over time
- ▶ Gating units: regulate the information flow into and out of the memory

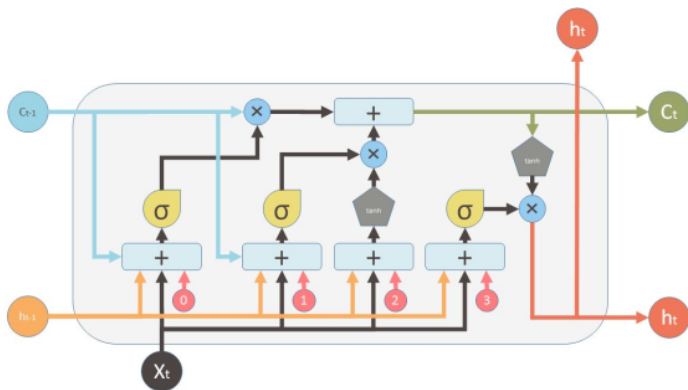


LSTM Memory Cell



Gate (sigmoid layer followed by pointwise multiplication)

LSTM Memory Cell



Inputs:



Input vector



Memory from previous block



Output of previous block

outputs:



Memory from current block



Output of current block

Nonlinearities:



Sigmoid



Hyperbolic tangent

Bias:



Vector operations:




Element-wise multiplication



Element-wise Summation / Concatenation

LSTM Cell state vector (C)

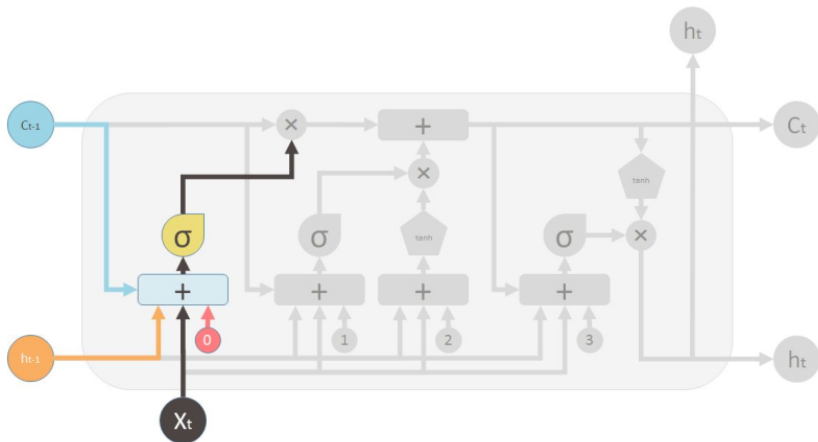
- ▶ Memory of the LSTM
- ▶ State can be changed by forgetting (\times) and addition of new data ($+$)
- ▶ Linear changes



LSTM Forget Gate

- Controls what remains of the previous memory

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$

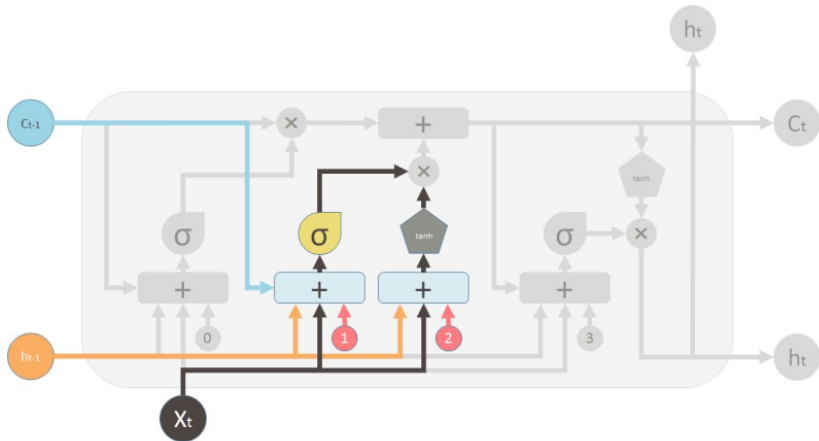


LSTM Input Gate

- Controls what new information is added to the memory

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$


$$\tilde{C}_t = \tanh(W_C x_t + U_C h_{t-1} + b_C)$$



LSTM Memory update

► Aggregation

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$




LSTM Output gate


- Conditionally decides what to output from the memory

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{h}_t = o_t * \tanh(C_t)$$



LSTM Memory Cell Summary



$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$


$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$

$$\tilde{C}_t = \tanh(W_C x_t + U_C h_{t-1} + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{h}_t = o_t * \tanh(C_t)$$



LSTM Training

- ▶ Number of parameters:
 - ▶ n number of LSTM units
 - ▶ m parameters in the input data
 - ▶ Dimension of U is $n \times m$
 - ▶ Dimension of W is $n \times n$
 - ▶ Dimension of b is n
 - ▶ There are four gates in an LSTM cell

$$\text{number of parameters} = 4(nm + n^2 + n)$$

$$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$$

$$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$$

$$\tilde{C}_t = \tanh(W_C x_t + U_C h_{t-1} + b_C)$$

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

$$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$$

$$\tilde{h}_t = o_t * \tanh(C_t)$$

LSTM Training

- ▶ Backpropagation Through Time (BPTT) most common
- ▶ Weights: Gates, input tanh layer
- ▶ Output:
 - ▶ One output at each timestep
 - ▶ Single output for the whole task

